

Level 1 → Level 2: the human **moves**

Same five stations. Count who stands at each one.

SET GOAL

START TASK

BUILD

CHECK

MERGE

LEVEL 1

human at 4 of 5



LEVEL 2

human at 2 of 5



the middle runs unattended

The human stops watching steps and starts designing the system.

Level definitions: Frontier Build NYC invite (Obvious). Station mapping is a schematic, not Obvious's product docs.

Seven stations of a software factory

0 · TRIGGER – a ticket, an event or a schedule starts the work.
No one types a prompt.

1	Spec Blueprint plus checklist: what "done" means	HUMAN
2	Agent plan Splits the spec into subtasks	MACHINE
3	Sandbox A prototype in a warehouse, not your kitchen	MACHINE
4	Multi-PR Change orders, each awaiting a signature	MACHINE
5	Agentic reviewer A second agent inspects each change	MACHINE
6	CI Test conveyor belt: did anything break?	MACHINE
7	Human merge Final walk-through into the real codebase	HUMAN


Order synthesized from the invite's Level 2 components, not
Obvious's published pipeline.

Six weeks of convergence

DATE	FACTORY	GITHUB	COGNITION
Aug 3	Agents as org actors		
Aug 19	\$100M "software factory" partners		
Aug 24	Routing belongs in the harness		
Aug 27	Builder / checker / judge roles		
Sept 1		AI approval counts toward merge	
Sept 8		Managed sandbox (JetBrains)	\$2B raised at \$48B
Sept 9	On Claude Marketplace	Agent permissions + autofix	
Sept 10			SWE-2, its own model
Sept 11		Review auto-resolution	Fusion harness

Pattern: harness as moat · checking as its own layer · sandbox + permissions first. Sept 17: Obvious hosts Frontier Build NYC.

Same model. Different jobs.



At Level 2, who writes the answer key?

