

Seat or outcome: what are you buying?

SOFTWARE (SaaS)

YOU BUY

A tool + a seat

PRICED

Per seat

WHO DOES THE WORK

Your team

WHO OWNS A BAD RESULT

You

SERVICES-AS-SOFTWARE

YOU BUY

The finished result

PRICED

Per outcome

WHO DOES THE WORK

**AI agents, checked
by experts**

WHO OWNS A BAD RESULT

The vendor

One room, six companies, one trade

	WHAT IT REPLACES	WHO CHECKS THE WORK	WHO BUYS
Amperos Health	Billing staff chasing insurers by phone; the scripted-bot first wave	Billing specialists take the edge cases	Hospital & clinic finance
Solstice Health	A ~3-month agency → brand → review-committee loop	Its own reviewers, before the committee	Pharma brand teams
Holly	HR staff hand-researching peer pay, months per role	Agency HR reviews the draft	Public-agency HR
Soxton	Hourly billing for routine founder paperwork	A licensed lawyer, every draft	Pre-launch founders
Warp	Hand-filed payroll, multi-state tax registrations and compliance notices	Warp Agent runs it; human review not disclosed	Startups & growing companies
Pascal	Retailer-portal exports re-keyed into spreadsheets	Not a regulated service	Consumer-goods brands

Sources: AlleyWatch; Becker's Hospital Review; TechFundingNews; Harvard Innovation Labs; BusinessWire; Pascal (company page); Warp (warp.co). 2023–26.

One more client, one more hire?

TRADITIONAL SERVICE FIRM



SERVICES-AS-SOFTWARE



What doesn't flatten: liability. Someone still owns the error.

Staffing reduced, not eliminated

~1/2

the staffing a traditional integrator uses: **Humanize**, implementing SAP, Salesforce and ServiceNow

70%

of \$18B+ in annual loan factoring automated: **KUNGFU.AI** for Triumph Financial

Both are 2026 Services-as-Software Award winners.
Note the verb: reduced.

Source: HFS Research, 2026 Services-as-Software Award winners.

Three tests

- 1 Is it priced per result, or per seat with a new name?
- 2 When the AI is wrong, who signs for it?
- 3 As volume grows, does review time shrink, or scale with it?

Which of the six passes all three?