

Cheaper tokens. Bigger bills.

price per token

total AI spend



Token Economics · VentureBeat salon · NYC · Sept 29

Price per token ranks last


Integration 41%



Total cost of ownership 35%



Cost per token 8%



Share of enterprises naming each as their top infrastructure-buying factor.

Source: VentureBeat Pulse Research, “The AI compute gap,” 2026.

THE MECHANISM

Where the bill actually comes from

1 · PROMPT

Read all at once



2 · KV CACHE

Sticky notes of every token read



chat



long context



agent chaining
many calls



3 · GPU MEMORY (HBM)

Small, scarce, pricey



4 · OUTPUT

Written one token at a time

More context + more agent steps = more notes held in the most expensive memory.

THREE VANTAGE POINTS

Three levers on the same bill

ARCHITECTURE

ACTS ON

Routing each sub-task to the right model or agent

IN THE ROOM

Mount Sinai · Benjamin Glicksberg

EVIDENCE

Published multi-agent clinical study: accuracy held under heavy load with **up to 65x fewer** compute resources

MEMORY

ACTS ON

Extending GPU memory so the KV cache fits

IN THE ROOM

WEKA · Alan McSeveney, Field CTO

MEASUREMENT

ACTS ON

Knowing which workflow costs what per useful result

IN THE ROOM

Venmo · Byron Ma, CTO

OPEN QUESTION

What does that look like at transaction scale?

Source: Mount Sinai / npj Health Systems, 2026 (65x figure).

Pricing problem, or measurement problem?

Pricing →

Buy cheaper capacity or more memory.

Measurement →

Give someone ownership of cost per useful result.

Which one is your team
fighting?