Speech-to-text had amnesia.

For as long as voice agents have existed, the layer that turns your voice into text has heard a chunk of audio, transcribed it, and instantly forgotten. This release finally gives it memory.

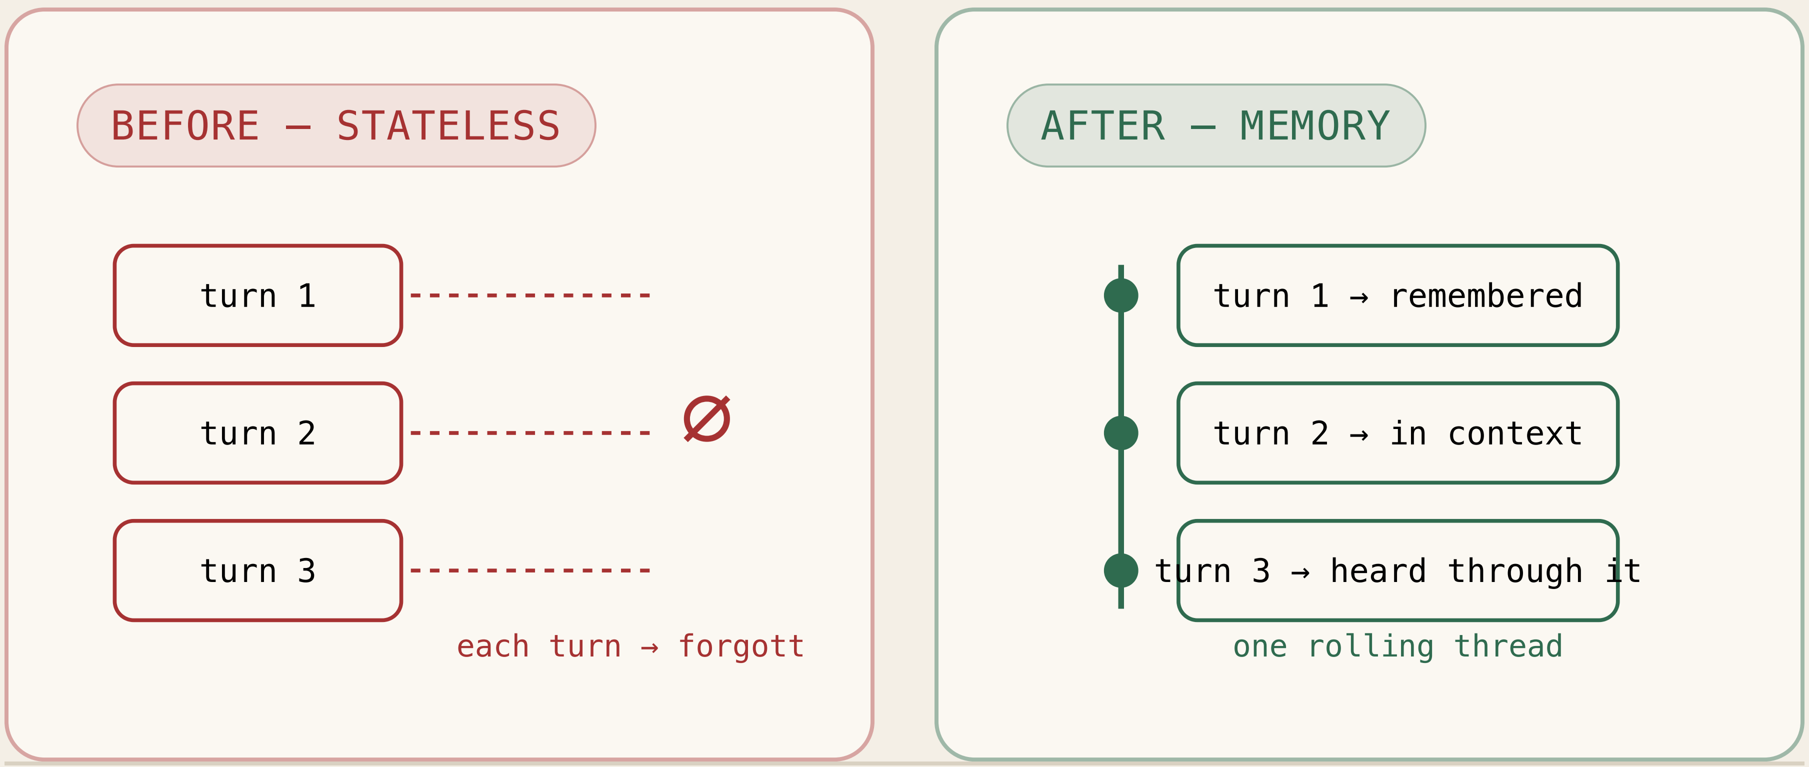

`agent_context`, visualized

The agent feeds its *own last question* back into transcription — so your reply is heard through the lens of what was asked. Live, **mid-call, no reconnect.**

WITHOUT — reply transcribed in isolation

agent asks → *“What’s your email?”*

caller says → “user at gmail dot com” → **5 random words · wrong action ✗**

WITH agent_context — heard through the question

context → an email is expected

caller says → “user at gmail dot com” → **user@gmail.com** ✓

The hard part isn’t the idea — it’s doing it **without tearing down the session.** Most streaming stacks can’t change context without dropping the connection. This one updates over the open socket.

STT = speech-to-text

Vendor number vs. independent number

Two sources, labeled — because a launch claim and a third-party benchmark are not the same evidence. Here, both point the same way.

VENDOR – AssemblyAI

−10.2%

word-error-rate, across 20,000 agent files

fabrications	**−18.3%**
hallucinations	**−17.2%**
place names	**−15.5%**

INDEPENDENT – Coval

3.40%

word-error-rate · **#1 of 29 models**

open Apache-2.0 harness
897 messy conversational clips
inside the 2–4% human-parity band

A model transcribing live, interrupted speech **as accurately as a professional human.**

Accuracy leader — not the speed leader

The accuracy is bought with latency. The tradeoff didn’t dissolve; it got repriced.

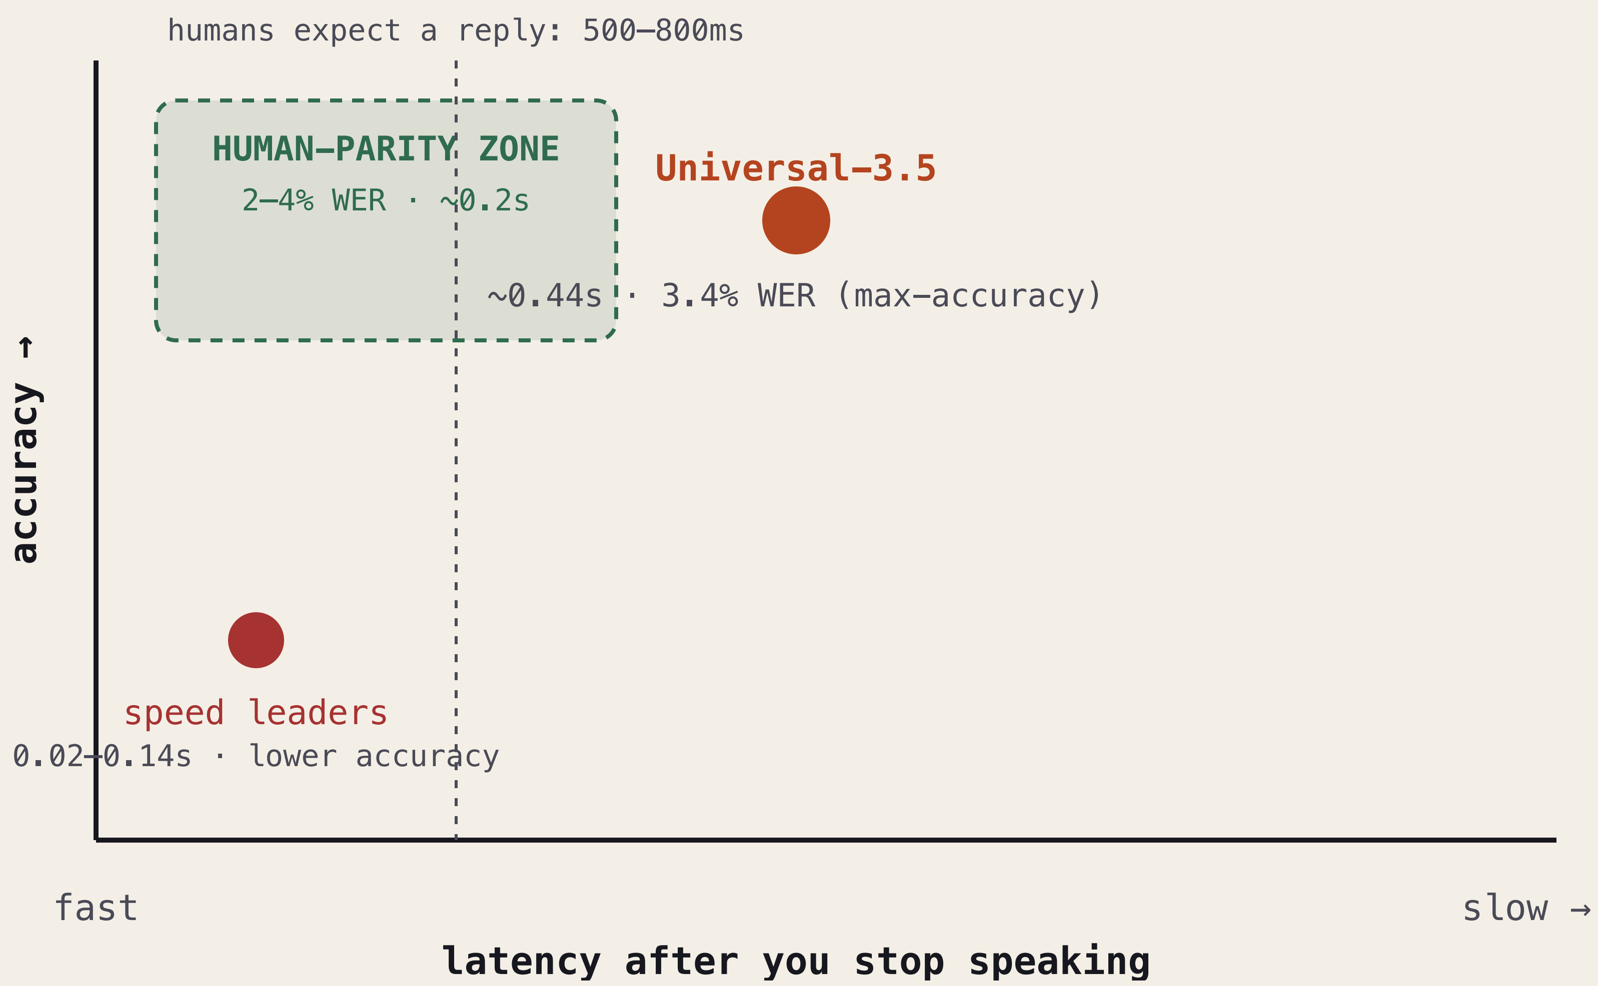

And **$0.45/hour is the base** — contextual prompting, diarization-with-revision, and voice isolation are add-ons that stack on top.

It got better at *what* it hears — not *when a turn ends*.

Knowing you're actually done, not just pausing, is **turn detection** — still the hardest unsolved problem in voice. Two architectures are betting against each other:

PATH A · pipeline

Context-accurate STT + a separate turn-detector

AssemblyAI's bet: win transcription accuracy, bolt on endpointing.

PATH B · fused

One model that natively knows when you're done

Deepgram Flux's bet: fold turn-detection into the transcript stream.

THE OPEN QUESTION

Nobody's proven you get both accuracy *and* native turn-detection yet.

April: network quality → June: small models → Sept: context at human parity → next wall: the turn.

which side are you on?