


The average hides the crossover.

GLM-5.3 is the #1 open-weight model — and that's the least interesting thing about it.

One score says: still a tier back.



#1 open-weight, ~8 points behind the closed frontier. Read the blended "how-smart" index alone, and open weights look a step behind.

On the work that ships agents, it crosses in.

LONG-HORIZON CODING · TERMINAL-BENCH V4

GLM-5.3		41.9%
GLM-5.3 Flash		32.8%

Top open-weight — and **ahead of a closed frontier coder** (GPT-5.6 Sol), per Artificial Analysis.

AGENTIC KNOWLEDGE WORK · AA-BRIEFCASE (ELO)

GLM-5.2		1234
GLM-5.3		1515
GLM-5.3 Flash		1455

+281 Elo over 5.2 in one release — shoulder-to-shoulder with the closed leaders.

WHY IT'S REAL, NOT A FLUKE

It wins by not overthinking.

Overthinking
model

too many turns → burns the budget → gets
stuck

GLM-5.3

token-efficient · thinking dialed in → stays on
task → finishes

"For models thinking way too much, trying to do too many turns — they get stuck... it's more token-efficient, its thinking levels are dialed in."

— Declan Jackson, Artificial Analysis

Sustained execution over a long horizon is a **different skill** than one-shot test IQ. And Flash rides the same strength to 3rd open-weight on both — so it's the pattern, not the flagship.

THE TAKEAWAY

When someone hands you one number for a model, ask what it's the average of.

On agentic coding + knowledge work, the open, self-hostable, fine-tunable model is **already at the frontier** — at open-weight economics.

\$0.87 vs ~\$6

same "build the deck from my files" task — Flash vs flagship

MIT · self-host · tune

Flash is MIT-licensed; both run on your own infra

Thanks to **Philip Kiely** (Baseten) & **Declan Jackson** (Artificial Analysis).