

Your org's build line, run by agents

- ① **Connect your repos** → prompt what you want built
↓
- ② An **orchestrator** plans a DAG of tasks (it runs over days)
↓
- ③ Each task → a specialized **agent in its own isolated sandbox**
↓
- ④ Inside: clone → **snapshot** → build → must pass QA to finish
↓
- ⑤ Work converges into a **PR** — QA · security · adversarial review → you approve

Every sandbox is **isolated and snapshotted** — so runs are reproducible, and you pay for compute only while agents are actually building.

One PR she showed ran ~17,000 lines – already reviewed by QA + security agents before it reached her inbox.