

How GRPO teaches a model to code

- 1 Give the model a **prompt** (a task you can score)
↓
- 2 Sample a whole **group** of attempts
↓
- 3 **Score** each one — did the code pass the tests?
↓
- 4 **Advantage** = each attempt vs. the group average
↓
- 5 **Nudge** the model toward the winners → repeat

The trick: **no critic model**. The group's own average is the baseline — so GRPO drops the second network PPO needs, which is why large-scale RL got affordable.

Live demo: [Qwen 2.5 0.5B](#) · [38 → 77](#) · [Daytona GPU sandbox](#)