

What enterprises actually buy AI infrastructure on


Integration 41%



Total cost of ownership 35%



Cost per token (ranks last) 8%



Share naming each as their top infrastructure-buying factor. **TCO** = the full cost of running it, not just the per-call price.

44%

rigorously track compute cost or ROI. The rest are managing a bill they can't fully see.

Source: VentureBeat Pulse Research, "The AI compute gap" (Infrastructure & Compute, June 2026).

Why cheaper tokens still mean bigger bills

1 · PROMPT

Read all at once



2 · KV CACHE

Sticky notes on every token already read



chat



long context



agent chaining
many calls



3 · GPU MEMORY (HBM)

Small, scarce, the priciest in the data center



4 · OUTPUT

Written one token at a time


Price per token falls; notes per task climb. That's the Jevons paradox: efficiency raises total use.

Mechanism: vLLM / PagedAttention literature; KV-cache explainers, 2026.

Open-weight models took most of the traffic in eight months


December 2025

<10%



April 2026

13%



August 2026

56%



Open-weight share of token volume on Vercel's AI Gateway. **Open-weight** = weights public, so anyone can host the model.

BLENDED PRICE, PER MILLION
TOKENS

\$18.40 → \$6.07

↓ 67%, Q1'25 → Q1'26

AVERAGE ENTERPRISE AI BUDGET

\$1.2M → \$7M

↑ 2024 → 2026

Sources: Vercel AI Gateway Production Index, Sept 2026; Correlation One, 2026 enterprise token-cost playbook.

One request, four seats on the panel

INGEST

Messy PDFs and scans → structured data for agents

EXTEND



ROUTE

The gateway picks a model for each request

VERCEL AI GATEWAY



SMALL / OPEN-WEIGHT

Easy work, cheap

FRONTIER / CLOSED

Hard work, expensive

– – escalate to the frontier model if unsure (a “cascade”) – –



SERVE

The inference host runs the models

BASETEN

PAY + GOVERN

The buyer: Brex, a Capital One business since April 2026

Roles per event invite and company sites; Capital One newsroom, Apr 2026.

Same questions, two paths, two scores

Claude Sonnet 4.6

Text-to-SQL

90%

Semantic layer

98.2%

GPT-5.3 Codex

Text-to-SQL

84.1%

Semantic layer

100%

31%

best model's score on Spider 2.0, a text-to-SQL test built on realistic enterprise databases.

Sources: dbt Labs, Semantic Layer vs. Text-to-SQL benchmark, Apr 2026; Spider 2.0 via arXiv 2507.14372 (2025).

Where a wrong answer goes

TEXT-TO-SQL

Plain-English question



AI writes a new query



Wrong join



Confident number,
no alert

SEMANTIC LAYER

Plain-English question



AI picks from approved metrics



Metric not defined



⚠ **Error:**
can't answer

Gartner expects 60% of agentic analytics projects relying only on MCP to fail by 2028 without a semantic layer.

MCP = the standard plug that lets an agent call tools; it says how to ask, not what a metric means.

Sources: dbt Labs, Apr 2026; Gartner Market Guide for Agentic Analytics, Feb 2026 (via Cube).